

Cleaning and Transforming Data

Feature engineering and transparent preparation



Prof. Gabriel Spadon, PhD

Faculty of Computer Science
spadon@dal.ca

Learning objectives and how to use this handout

By the end of this class, students should be able to

- name the three missing-data mechanisms and state what each allows the missingness to depend on;
- replace a missing value in a separate copy while flagging it and keeping the source table;
- build a derived variable with correct units, and standardize a numeric variable with a stated divisor;
- draw a random and a stratified sample, and compare equal-width and equal-frequency bins on the same data;
- encode a nominal variable with indicators rather than an invented order; and
- turn a monitoring record into a documented table that separates exact values, reporting limits, missing values and invalid tokens.

Each section has the same four parts. It says what the operation is for, gives the rule or formula, works one small example, and shows the R and Python code that produces the numbers. The sections get harder as you go. Sections 1–7 each use one operation on a small vector. Section 8 combines them on a monitoring record. The two optional extensions (model-based sampling and principal component analysis, PCA) are not required for Lab 02.

1 Missing data

What it is for Before choosing how to handle a blank, you need an assumption about why it is blank.

The rule A mechanism describes the probability that a value is unavailable. It is a property of how the data were collected and recorded, so one blank cell cannot identify it. Table 1 states what the missingness may depend on under each mechanism.

Table 1. What missingness can depend on under each mechanism.

Mechanism	What missingness can depend on
MCAR	Neither observed nor unobserved values.
MAR	Observed values; after conditioning on them, not the missing values.
MNAR	Unobserved values still matter after conditioning on observed information.

- **MCAR.** The complete cases are a random subset of the eligible records. Complete-case summaries avoid selection bias from missingness, but they are less precise.
- **MAR.** Missingness may vary with observed information. An analysis can account for that information, but MAR alone does not make an arbitrary imputation or deletion rule valid.
- **MNAR.** Dependence on unobserved values remains after conditioning on what is observed. Handling it needs external information, an explicit model, or a sensitivity analysis [9, 13].

A failure of MCAR does not mean every complete-case estimate is biased. Some correctly specified conditional models stay valid under broader conditions. Ask which quantity you are estimating and which assumption supports it.

The strategies Figure 1 lists the alternatives. The branches are choices, not steps to run in order.




Figure 1. Alternative missing-data strategies. Branches identify alternatives, not a required sequence of operations.

- **Listwise deletion** keeps only records complete on the variables an analysis needs. **Pairwise deletion** uses the available records separately for each statistic, so a mean for one variable can use rows where another variable is missing. Report the denominator of every summary.
- **Mean, median and mode replacement** are single imputation. Filling with the observed mean leaves the observed mean unchanged, which is arithmetic and says nothing about the population mean. Adding copies of the mean adds zero deviations to the numerator of the variance while the divisor grows, so the sample variance shrinks. Median or mode replacement need not preserve the mean. None of these restores the uncertainty of the missing measurements.
- **LOCF** (last observation carried forward) repeats the previous value in a series. **BOCF** (baseline observation carried forward) repeats a designated baseline value. Backward fill uses the next observed value and is not BOCF. Sort by time and keep units, people or sites separate before any carry. A leading missing value has no earlier value to carry.
- **Regression imputation** predicts a missing value from other variables. The mean prediction alone omits residual variation. **Multiple imputation** builds several plausible completed datasets, analyses each, and pools the estimates and their uncertainty. The `mice` package implements it [3]. These model-based methods are not required for Lab 02.

Worked example A column holds 2, a missing value, 8 and 12. The observed values are 2, 8 and 12, so the observed median is 8. The code below records which row was missing, keeps the source table, and fills the missing value with 8 in a separate copy.

R

```
raw <- data.frame(value = c(2, NA, 8, 12))
observed <- raw[!is.na(raw$value), ,
                drop = FALSE]

cleaned <- raw
centre <- median(observed$value)
cleaned$was_missing <- is.na(raw$value)
cleaned$value[cleaned$was_missing] <- centre
cleaned
# An illustrative choice, not a recovered
# measurement.
```

Python

```
import pandas as pd

raw = pd.DataFrame({"value": [2, None, 8, 12]})
observed = raw.dropna(subset=["value"])
cleaned = raw.copy()
centre = observed["value"].median()
cleaned["was_missing"] = raw["value"].isna()
cleaned["value"] = (
    cleaned["value"].fillna(centre))
print(cleaned)
```

The flag and the untouched source are the evidence for the decision. The example does not justify median replacement for a monitoring dataset. In Lab 02, keeping a missing value and reporting its count is often the most defensible choice.

2 What data preparation does

What it is for Data preparation is the CRISP-DM stage that turns the table you received into the table an analysis needs. Cleaning, transforming, feature engineering, and reduction are related preparation tasks.

The rule Preparation is iterative. A summary may reveal an unexpected unit, a plot may show a recording error, and either can send you back to change an import rule. You do not have to finish every preparation step before looking at the data again. Table 2 pairs four common problems with the operation that answers each. The rest of this handout takes them one at a time.

Table 2. Data problems and possible preparation responses.

Problem	Possible response
Variables have incompatible scales	Data transformation
A useful ratio needs its own variable	Feature engineering
Records contain unavailable or invalid values	Data cleaning
Too many variables obscure an analysis	Dimensionality reduction

3 Feature engineering and units

What it is for A derived variable combines source measurements into a quantity that answers the question better than either source alone.

The rule Write the formula with its units and convert every input to those units before computing. Body mass index (BMI) is

$$\text{BMI} = \frac{\text{mass in kg}}{(\text{height in m})^2}.$$

Worked example For a mass of 82 kg and a height of 178 cm, convert the height to 1.78 m first. Then

$$\text{BMI} = \frac{82}{1.78^2} = 25.88 \text{ kg/m}^2.$$

The code applies the same formula to three people, giving 25.88, 22.41 and 29.00. BMI here illustrates arithmetic and documentation, not a clinical judgement.

R

```
measurements <- data.frame(
  mass_kg = c(82, 61, 95),
  height_cm = c(178, 165, 181))
measurements$bmi <- measurements$mass_kg /
  (measurements$height_cm / 100)^2
round(measurements$bmi, 2)
# 25.88 22.41 29.00
```

Python

```
import pandas as pd

measurements = pd.DataFrame({
  "mass_kg": [82, 61, 95],
  "height_cm": [178, 165, 181]})
measurements["bmi"] = (
  measurements["mass_kg"]
  / (measurements["height_cm"] / 100)**2)
print(measurements["bmi"].round(2).to_list())
# [25.88, 22.41, 29.0]
```

Both versions keep the original measurements beside the derived value. The columns come from one table, so corresponding rows describe the same person. R combines vectors by position and can recycle a shorter vector. pandas combines Series by their index labels. When the inputs come from different tables, establish the row relationship first. Equal lengths do not prove that rows correspond.

4 Reducing a dataset and sampling

What it is for Reduction makes a dataset smaller along one of three axes. Each discards something, so choose it for an analytical reason, not because the file gets smaller.

The three choices Figure 2 separates them.

- **Fewer features.** Simplifies computation or allows a two-dimensional display. Feature selection keeps some original variables. Feature projection builds new coordinates from combinations of variables. PCA is a linear projection, developed in the optional extension.
- **Fewer rows.** Makes a first inspection manageable, but may omit rare conditions.
- **Fewer distinct values.** Binning turns numbers into categories and loses the detail inside each bin (Section 6).




Figure 2. Three separate reduction choices. PCA is developed in the optional extension.

Random and stratified samples A random sample chooses records through a defined chance mechanism. A stratified sample first splits the records into groups, then samples within each group. Stratification helps keep rare groups, and the grouping variable and allocation must follow the question. It cannot repair a biased source dataset.

Worked example Ten records come from two sites, five from each. Both samples below contain six records. The random sample may take any mix of sites. The stratified sample takes three from each site, so its site proportions (50 percent each) match the source table.

R

```
data <- data.frame(row_id = 1:10,
  site = rep(c("A", "B"), each = 5))
set.seed(3141)
idx <- sample(seq_len(nrow(data)), size = 6)
random_sample <- data[idx, ]
# Draw three rows from each site.
idx_by_site <- split(
  seq_len(nrow(data)), data$site)
chosen <- unlist(
  lapply(idx_by_site, sample, size = 3))
stratified_sample <- data[chosen, ]
table(stratified_sample$site)
# A 3, B 3
```

Python

```
import pandas as pd

data = pd.DataFrame({
  "row_id": range(1, 11),
  "site": ["A"]*5 + ["B"]*5})
random_sample = data.sample(
  n=6, random_state=3141)
# Draw three rows from each site.
stratified_sample = data.groupby("site").sample(
  n=3, random_state=3141)
print(stratified_sample["site"].value_counts())
# A 3, B 3
```

Differences to know

- A seed makes a draw reproducible within one software and algorithm. `set.seed(3141)` and `random_state=3141` do not give the same rows.
- Pass an integer sample size. Seventy percent of 101 rows gives 70 rows after R's integer truncation, while `pandas sample(frac=0.7)` rounds to 71. Sampling a fraction within each group can change the total again.

With proportional allocation, every record has the same chance of selection. If groups are sampled at different rates, a population estimate must weight by those probabilities. Records from the same station or period may be dependent, so randomly chosen rows are not automatically independent observations. Lab 02 does not require subsampling or a training/test split.

5 Standardization

What it is for A variable's units set its numerical spread. Methods that compare variables through distances or variance let the largest scale dominate. Standardization puts every variable on a unit-free scale.

The formula For observed values with a nonzero sample standard deviation,

$$z_i = \frac{x_i - \bar{x}}{s_x},$$

where $\bar{x}$ is the observed mean and s_x is the sample standard deviation with divisor $n - 1$. The results have no units, mean zero and sample standard deviation one.

Worked example For the vector $x = 2, 4, 4, 5, 7, 9, 12, 41$, the mean is $\bar{x} = 10.5$ and $s_x = 12.72792$. The first value becomes $z_1 = (2 - 10.5)/12.72792 = -0.6678$. Table 3 gives selected values.

Table 3. Selected standardized values, using the sample standard deviation.

x_i	$x_i - \bar{x}$	z_i
2	-8.5	-0.668
4	-6.5	-0.511
7	-3.5	-0.275
12	1.5	0.118
41	30.5	2.396

R

```
x <- c(2, 4, 4, 5, 7, 9, 12, 41)
z <- scale(x)
attr(z, "scaled:center") # 10.5
attr(z, "scaled:scale") # 12.72792
round(z[1], 4) # -0.6678
scale(c(x, NA))[1] # -0.6678231
# Only the last output is NA; observed
# values remain.
```

Python

```
import numpy as np
from sklearn.preprocessing import StandardScaler

x = np.array([2, 4, 4, 5, 7, 9, 12, 41.])
s = StandardScaler().fit(x.reshape(-1, 1))
print(s.mean_[0]) # 10.5
print(round(s.scale_[0], 4)) # 11.9059
print(round(s.transform(x[:, None])[0, 0], 4))
# -0.7139; use ddof=1 for the R convention:
print(round(
    ((x - x.mean()) / x.std(ddof=1))[0], 4))
# -0.6678
```

Differences to know

- R's `scale()` divides by the sample standard deviation (divisor $n - 1$) [12]. scikit-learn's `StandardScaler` uses divisor n [11]. For this vector its standard deviation is 11.9059, so the first score is -0.7139 instead of -0.6678. Use `std(ddof=1)` to reproduce R.
- Both ignore missing values when estimating centre and spread, and both keep missing positions in the output.
- A constant or all-missing column needs an explicit decision. Division by zero is not a standardization. `StandardScaler` gives a constant column a scale of one, which does not give it variance one.

Standardization changes the scale, not the shape. It does not remove outliers, skewness or dependence. The value 41 is still the extreme observation, now 2.396 sample standard deviations above the mean. A one-unit change in z is a numerical comparison, not proof that two variables matter equally.

6 Discretization

What it is for Discretization replaces a number with an interval label, for a question that is naturally about categories. It discards within-bin differences and makes the result depend on the chosen boundaries, so it is not automatic cleaning.

The two rules

- **Equal width.** k bins whose boundaries are evenly spaced in measurement units, $(\max - \min)/k$ apart.
- **Equal frequency.** Boundaries at sample quantiles, so each bin holds about the same number of values. Ties make the counts approximate.

The first interval includes its lower boundary. Each later interval is open on the left and closed on the right.

Worked example Use four bins on the same vector $x = 2, 4, 4, 5, 7, 9, 12, 41$. Equal-width spacing is $(41 - 2)/4 = 9.75$, giving edges 2, 11.75, 21.5, 31.25, 41 (Table 4). The sample quartiles give equal-frequency edges 2, 4, 6, 9.75, 41 (Table 5).

Table 4. Equal-width bins with explicit boundaries.

Interval	Count	Members
[2, 11.75]	6	2, 4, 4, 5, 7, 9
(11.75, 21.5]	1	12
(21.5, 31.25]	0	none
(31.25, 41]	1	41

Table 5. Equal-frequency bins based on sample quartiles.

Interval	Count	Members
[2, 4]	3	2, 4, 4
(4, 6]	1	5
(6, 9.75]	2	7, 9
(9.75, 41]	2	12, 41

The counts are 6, 1, 0, 1 for equal width and 3, 1, 2, 2 for equal frequency. Equal width keeps equal spacing between boundaries, but six of eight values share one bin. Equal frequency spreads the counts, but its last bin holds both 12 and 41 and hides their 29-unit difference.

R

```
x <- c(2, 4, 4, 5, 7, 9, 12, 41)
edges <- seq(min(x), max(x), length.out = 5)
width_bins <- cut(x, breaks = edges,
                  include.lowest = TRUE)
quantiles <- quantile(
  x, probs = seq(0, 1, 0.25))
stopifnot(length(unique(quantiles)) == 5)
freq_bins <- cut(x, breaks = quantiles,
                include.lowest = TRUE)
table(width_bins) # 6 1 0 1
table(freq_bins) # 3 1 2 2
```

Python

```
import numpy as np
import pandas as pd

x = np.array([2, 4, 4, 5, 7, 9, 12, 41])
edges = np.linspace(x.min(), x.max(), 5)
width_bins = pd.cut(
  x, bins=edges, include_lowest=True)
freq_bins = pd.qcut(x, q=4)
print(pd.Series(width_bins).value_counts(
  sort=False))
print(pd.Series(freq_bins).value_counts(
  sort=False))
```

Differences to know

- Given an integer number of bins, both cut functions can extend the range slightly to include the endpoints. With explicit edges, set `include.lowest = TRUE` in R or `include_lowest=True` in pandas so the minimum falls in the first bin.
- With many ties, quantile boundaries can repeat. R's `cut` rejects non-unique breaks, and `qcut` raises by default.

When boundaries repeat, use fewer intervals, choose meaningful boundaries, or keep the continuous variable. Do not split identical values across categories. Report the number of bins you obtained, not the number you asked for.

7 Representing categories

What it is for Some software needs numbers, not labels. The representation must not add information the labels do not have.

The rule Encode a nominal variable with one indicator column per category. A row has 1 in the column of its category and 0 elsewhere. Coding sites A, B and C as 1, 2 and 3 would claim an order and equal spacing that the labels do not have.

Worked example The site labels A, B, A, C give three indicator columns. The first and third rows have 1 under A, the second under B and the fourth under C.

R

```
site <- factor(c("A", "B", "A", "C"))
indicators <- model.matrix(~ site - 1)
indicators
# Three columns, one per site; no intercept.
```

Python

```
import pandas as pd

site = pd.Series(["A", "B", "A", "C"])
indicators = pd.get_dummies(site, dtype=int)
print(indicators)
# Three columns, one for each site.
```

Neither version assigns an order to the site names. Lab 02 needs clear variable types and documented labels. It does not need indicator encoding or a model.

8 From monitoring records to a documented table

What it is for This section combines the operations above on one monitoring record, the kind of file Lab 02 uses. The goal is an analytical table in which every change can be traced back to the raw record.

8.1 Separate the meaning of unavailable measurements

The rule Read every field as text first, list the distinct tokens, and give each one a status before any numeric conversion. A value that cannot be used in exact arithmetic may still carry information.

Worked example An illustrative concentration column is documented in mg/L. Table 6 shows five strings with four different statuses.

Table 6. A small monitoring example with distinct value-status categories.

Raw field	Status	Interpretation
0	measured	Exact recorded zero in the stated unit.
blank	missing	No usable measurement was recorded.
<0.1	censored	Below 0.1 mg/L; retain 0.1 as a limit, not a value.
oops	invalid	Unexpected token; investigate the original record.
0.24	measured	Exact recorded value, subject to provider quality flags.

- 0 and 0.24 are exact measurements. A measured zero is a value.
- <0.1 is censored. It reports an upper limit, not a measurement of 0.1 and not a zero.
- The blank is missing under this example's documentation.

- oops is invalid under the numeric schema and needs investigation.

Provider flags may add restrictions, such as a provisional reading or a failed instrument check. The code below assigns the statuses by hand, keeps the raw strings and units, and checks that the raw text is unchanged.

R

```
raw <- data.frame(
  value_text = c("0", "", "<0.1", "oops",
                "0.24"),
  unit = rep("mg/L", 5),
  stringsAsFactors = FALSE)
cleaned <- raw
cleaned$status <- c(
  "measured", "missing", "censored",
  "invalid", "measured")
cleaned$value_mg_L <- c(0, NA, NA, NA, 0.24)
cleaned$limit_mg_L <- c(NA, NA, 0.1, NA, NA)
stopifnot(identical(
  raw$value_text, cleaned$value_text))
# These statuses follow the stated rules.
# Original strings and units remain.
```

Python

```
import pandas as pd

raw = pd.DataFrame({
  "value_text":
    ["0", "", "<0.1", "oops", "0.24"],
  "unit": ["mg/L"]*5})
cleaned = raw.copy()
cleaned["status"] = [
  "measured", "missing", "censored",
  "invalid", "measured"]
cleaned["value_mg_L"] = [
  0, None, None, None, 0.24]
cleaned["limit_mg_L"] = [
  None, None, 0.1, None, None]
assert raw["value_text"].equals(
  cleaned["value_text"])
```

Hand assignment works for five known records. A real file needs a rule taken from the provider's data dictionary. Keep literal tokens such as NA as text until the documentation says whether they mean unavailable, a category, or something else. Do not suppress conversion warnings without looking at the records that caused them.

8.2 Apply the documented rule to every row

The rule For this example only, a blank means missing, an unsigned decimal string is an exact measurement, and a less-than sign followed by such a decimal is a reporting limit. Anything else is invalid and is not silently turned into missing. The regular expressions encode these rules.

- `[0-9]+` matches one or more digits.
- `([.][0-9]+)?` allows an optional decimal part.
- The anchors `^` and `$` in R, or the full-match operation in Python, require the whole string to match.

Worked example Applied to the five strings, the rule reproduces the statuses of Table 6. Exact values go to `value_mg_L`, the limit 0.1 goes to a separate `limit_mg_L` column, and the exact values are converted with $1 \text{ mg/L} = 1000 \mu\text{g/L}$, so 0.24 mg/L becomes $240 \mu\text{g/L}$.

R

```
raw <- data.frame(
  value_text = c("0", "", "<0.1", "oops",
                 "0.24"),
  unit = rep("mg/L", 5))
cleaned <- raw
text <- trimws(raw$value_text)
# This provider defines blank as missing.
missing <- text == ""
censored <- grepl("<[0-9]+([.][0-9]+)?$", text)
measured <- grepl("[0-9]+([.][0-9]+)?$", text)
cleaned$status <- ifelse(missing, "missing",
  ifelse(censored, "censored",
    ifelse(measured, "measured", "invalid")))
cleaned$value_mg_L <- NA_real_
cleaned$value_mg_L[measured] <-
  as.numeric(text[measured])
cleaned$limit_mg_L <- NA_real_
cleaned$limit_mg_L[censored] <-
  as.numeric(sub("<", "", text[censored]))
cleaned$value_ug_L <- 1000 * cleaned$value_mg_L
cleaned
```

Python

```
import pandas as pd
raw = pd.DataFrame({
  "value_text":
    ["0", "", "<0.1", "oops", "0.24"],
  "unit": ["mg/L"]*5})
cleaned = raw.copy()
text = raw["value_text"].str.strip()
# This provider defines blank as missing.
missing = text.eq("")
censored = text.str.fullmatch(
  r"<[0-9]+([.][0-9]+)?")
measured = text.str.fullmatch(
  r"[0-9]+([.][0-9]+)?")
cleaned["status"] = "invalid"
for mask, label in [(missing, "missing"),
                    (censored, "censored"),
                    (measured, "measured")]:
    cleaned.loc[mask, "status"] = label
cleaned["value_mg_L"] = pd.to_numeric(
  text.where(measured))
cleaned["limit_mg_L"] = pd.to_numeric(
  text.where(censored).str.removeprefix("<"))
cleaned["value_ug_L"] = (
  1000 * cleaned["value_mg_L"])
print(cleaned)
```

These rules are deliberately narrow. Real records may use signs, scientific notation, other missing codes, or other limit formats, so inspect the tokens before adapting them. The documentation and the token audit decide the classification, and the code makes it reproducible across all rows. The cleaned copy never replaces a censored result with zero or half the limit. That would need an extra assumption and a sensitivity analysis, and Helsel shows how such substitution distorts later estimates [7].

8.3 Check units, keys and counts

The rule Keep the downloaded file unchanged. Before combining observations, identify the station, time, parameter, unit and quality-flag fields, and state the observational unit. Many tables hold one row per station, time and parameter, not one row per station.

Duplicates The observational unit defines a duplicate. Choose the key columns, for example station, date and parameter, and count rows that share a key before removing anything.

- R. `duplicated(d[key]) | duplicated(d[key], fromLast = TRUE)` marks every row in a repeated group.
- Python. `d.duplicated(subset=key, keep=False)` does the same.

Worked example Four nitrate rows come from two stations. Station S1 has two rows on 2026-09-01, so the key station, date and parameter repeats. Both rows are marked, and neither is removed. The two values, 0.24 and 0.31 mg/L, differ, so the repeat is not an exact double export.

R

```
d <- data.frame(
  station = c("S1", "S1", "S1", "S2"),
  date = c("2026-09-01", "2026-09-01",
           "2026-09-02", "2026-09-01"),
  parameter = rep("nitrate", 4),
  value_mg_L = c(0.24, 0.31, 0.22, 0.18)
)
key <- c("station", "date", "parameter")
repeated <- duplicated(d[key]) |
  duplicated(d[key], fromLast = TRUE)
sum(repeated) # 2
d[repeated, ]
#   station      date parameter value_mg_L
# 1     S1 2026-09-01   nitrate      0.24
# 2     S1 2026-09-01   nitrate      0.31
```

Python

```
import pandas as pd

d = pd.DataFrame({
  "station": ["S1", "S1", "S1", "S2"],
  "date": ["2026-09-01", "2026-09-01",
           "2026-09-02", "2026-09-01"],
  "parameter": ["nitrate"] * 4,
  "value_mg_L": [0.24, 0.31, 0.22, 0.18],
})
key = ["station", "date", "parameter"]
repeated = d.duplicated(subset=key, keep=False)
print(repeated.sum()) # 2
print(d[repeated])
#   station      date parameter value_mg_L
# 0     S1 2026-09-01   nitrate      0.24
# 1     S1 2026-09-01   nitrate      0.31
```

A repeated key is a question for the documentation. Two rows may be an accidental double export, or they may differ in depth, sample fraction or laboratory replicate, which means the key was incomplete.

Units Convert only with a named rule, such as 1 mg/L = 1000 µg/L. Keep the original unit field and put the converted value in a new, clearly named column. Do not combine concentrations of different parameters because they share a unit. Flagging an unusual reading is a reason to investigate, not evidence that it should be deleted.

Counts Report the number of exact, missing, censored, invalid and excluded records. A mean of the exact measurements describes that subset. When censoring or exclusions are substantial it may not represent all sampled conditions, so state the denominator and the limitation beside the summary.

8.4 The Lab 02 workflow

Figure 3 and the list below give the order of work for one monitoring record.

1. Read one approved monitoring record and identify its observational unit.
2. Inspect types, units, missing tokens, reporting limits and quality flags.
3. Create a cleaned copy and record a reason for every changed or excluded value.
4. Compare summaries before and after preparation while keeping the raw record.
5. If categories help the question, compare two binning rules on one suitable numeric variable.

No PCA, model fitting or training/test split is required for Lab 02. Use one approved language. Follow the syllabus for the private submission and later portfolio release, and preserve the same submitted commit and its recorded history.

The worked examples in this handout are ungraded. The Lab 02 handout defines the assessed dataset, deliverables and submission.




Figure 3. A preparation workflow that keeps raw evidence separate from the analytical copy.

Integrated lab | Lab 02 | Cleaning water and air quality data

Build a transparent analytical table from one official monitoring record.

- **Carry forward.** Week 01 summaries, units, missingness, and unusual-value investigation, plus the reporting-limit rules from this class.
- **Start in class.** Choose one verified dataset, work through the selection checklist, and set up the R or Python analysis.
- **Assessment source.** week-02/Assessment/lab-02

9 Summary

- A missing-data mechanism is an assumption about the recording process. Any replacement goes in a flagged copy, never over the source.
- A derived variable is only as correct as its units. Convert before you compute.
- Standardizing changes scale, not shape. State the divisor, because R and scikit-learn differ.
- Equal-width and equal-frequency bins answer different questions, and both lose within-bin detail.
- Indicators encode membership without inventing an order.
- In a monitoring record, measured, missing, censored and invalid are different statuses. Keep the raw string, store limits apart from values, and report every count.

10 Optional extension on model-based sampling

These sections introduce model-evaluation vocabulary. They are not required for Lab 02.

10.1 Training, validation and test roles

What it is for When data are used to choose and evaluate a model, the rows that judge the model must not have shaped it.

The rule

- **Training** observations fit a model and its data-dependent transforms.
- **Validation** observations help choose among candidates or settings, such as sample size or feature count.
- **Test** observations support a final evaluation after those choices are fixed.

A fitted transform, such as a scaler or PCA, is fitted on the training rows and applied unchanged to the others. Fitting it on all rows before a comparison leaks information from the evaluation rows into training. This concerns model evaluation only. A descriptive summary of the current data needs no split. Related rows from the same station, person or period may need to stay together, because a random row split does not guarantee independence.

10.2 Why a simpler representation can help

Overfitting occurs when a model follows training-specific variation and loses accuracy on new observations. Fewer features can reduce computation and simplify interpretation, but can also discard useful information. A validation comparison shows when adding features stops helping (Figure 4). The test set must not be used repeatedly to choose the number of features.




Figure 4. A schematic comparison of training and validation error as more variables are kept. The curves are illustrative and do not determine an optimal feature count.

10.3 Model encodings

A linear model's intercept is its baseline prediction when the numeric inputs are zero. If every category has an indicator, the indicators sum to one, so including all of them with an intercept is redundant. One reference indicator is usually dropped. R's default treatment contrasts and an encoder with `drop="first"` do this. It is not a rule for every method. New categories at prediction time need an explicit policy, because encoding an unknown label as all zeros does not mean it was seen during fitting.

10.4 Incremental sampling

Incremental sampling starts with a small training sample and increases its size (Figure 5). Compare the candidates on the same validation set with a planned tolerance and a repeatability check. Judge a small change against the variation between repetitions, and keep relevant groups represented. One flat increment does not prove that more data cannot help.




Figure 5. An incremental sampling loop with a continuation decision and a stopping condition.

10.5 Combining predictions from several models

A method can be fitted separately to several samples, and each fitted model predicts the same new case (Figure 6). Numeric predictions are averaged. Class predictions are combined by voting. Bagging [2] fits each model to a bootstrap resample [4] of the training set, and those samples need not be smaller than the original.




Figure 6. Separate fitted models receive the same new case; their predictions are aggregated.

11 Optional extension on principal component analysis

11.1 Directions of variation

What it is for PCA replaces several correlated numeric variables with a few new coordinates that keep most of their variance [8, 10].

The rule Each principal component is a linear combination of the inputs. The first direction maximizes the projected variance among unit-length directions. Each later direction is perpendicular to the earlier ones and maximizes the remaining variance. If a two-variable cloud is long and narrow, one coordinate along its long direction approximates most of its variation (Figure 7). “Information” here means variance. A low-variance direction can still matter for a particular question or outcome.

Projection differs from selecting original columns. Computing a component score still needs every measurement with a nonzero weight.




Figure 7. The same illustrative cloud before and after adding PC1 and PC2 directions. Both arrows are drawn at the same length and do not encode variance. Both axes use the same physical unit scale.

11.2 Loadings on standardized iris measurements

Worked example The fit uses the four numeric iris measurements [5], collected by Edgar Anderson [1]. The species label is excluded. Each measurement is standardized so that their spreads are compared on one scale. Scaling is a choice. A covariance PCA suits cases where the original scales should be kept. Missing values and constant columns need a decision before this fit.

R

```
features <- iris[, 1:4]
p <- prcomp(features, scale. = TRUE)
# One component per column.
round(p$rotation, 4)
round(p$x[1, 1], 6) # -2.257141
round(p$sdev^2 / sum(p$sdev^2), 4)
# 0.7296 0.2285 0.0367 0.0052
```

Python

```
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

features = load_iris(as_frame=True).data
Z = StandardScaler().fit_transform(features)
p = PCA().fit(Z)
# One component per row.
print(p.components_.round(4))
print(p.explained_variance_ratio_.round(4))
# [0.7296, 0.2285, 0.0367, 0.0052]
```

Table 7. Loadings from sample-standardized iris measurements. Component signs follow the R example and may be reversed by another implementation.

	PC1	PC2	PC3	PC4
Sepal.Length	0.5210659	-0.3774176	0.7195664	0.2612863
Sepal.Width	-0.2693474	-0.9232957	-0.2443818	-0.1235096
Petal.Length	0.5804131	-0.0244916	-0.1421264	-0.8014492
Petal.Width	0.5648565	-0.0669420	-0.6342727	0.5235971

Let z_{SL} , z_{SW} , z_{PL} and z_{PW} be the centered, sample-standardized sepal length, sepal width, petal length and petal width. The first loading column of Table 7 gives

$$PC1 = 0.5210659z_{SL} - 0.2693474z_{SW} \\ + 0.5804131z_{PL} + 0.5648565z_{PW}.$$

The first iris observation scores -2.257141 on this component. Multiplying raw centimetre measurements by these weights gives a different, incorrect value. PC1 contrasts sepal width with the other three measurements. PC2 is dominated by sepal width.

Each loading column has unit length and the columns are mutually orthogonal. Component signs are arbitrary. Reversing a direction and all its scores describes the same projection.

Differences to know

- R stores one component per column in rotation. scikit-learn stores one per row in `components_`.
- R standardizes with divisor $n - 1$ and `StandardScaler` with n , so the input matrices differ by the factor $\sqrt{(n - 1)/n}$. Up to sign, R's scores equal Python's scores times that factor. Loadings and explained-variance proportions agree.
- `prcomp` uses the singular value decomposition. No manual eigenvector computation is needed.

11.3 Scores and a two-dimensional view

Multiplying the standardized data by the loading matrix gives the scores, one coordinate per observation and component. The score plot places each flower at its first two scores. Species shapes are added after fitting.

R

```
features <- iris[, 1:4]
p <- prcomp(features, scale. = TRUE)
png("iris-pca-R.png", width = 700,
    height = 700)
plot(p$x[, 1:2], pch = as.integer(iris$Species),
     xlab = "PC1", ylab = "PC2", asp = 1)
legend("topright",
     legend = levels(iris$Species),
     pch = 1:3, bty = "n")
dev.off()
```

Python

```
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
data = load_iris()
Z = StandardScaler().fit_transform(data.data)
scores = PCA().fit_transform(Z)
for k, marker in enumerate(["o", "^", "+"]):
    points = scores[data.target == k]
    plt.scatter(
        *points[:, :2].T, marker=marker,
        label=data.target_names[k])
plt.xlabel("PC1"); plt.ylabel("PC2")
plt.axis("equal"); plt.legend()
plt.savefig("iris-pca-Python.png")
plt.close()
```

A reflected axis between R and Python is an allowed sign choice. The visible grouping describes this dataset. It does not establish predictive accuracy on new observations.

11.4 Explained variance and retained components

The rule Each component's proportion is its variance divided by the total. Retain components until the cumulative proportion meets a target chosen for the intended use.




Figure 8. The first two sample-standardized iris scores. Species shapes are labels for interpretation, not inputs to PCA.

Worked example The proportions are 0.7296, 0.2285, 0.0367 and 0.0052. The cumulative proportions are 0.7296, 0.9581, 0.9948 and 1.0000. Two components retain 95.81 percent of the variance and exceed an illustrative 90 percent target.




Figure 9. Per-component variance on the left and cumulative retained variance on the right. The dashed line marks an illustrative 90 percent target.

The target is a choice, not a theorem, and it does not guarantee that the discarded variation is irrelevant. If the representation feeds a predictive model, choose the number of components within the training and validation process above.

11.5 Biplots and their display scales

A biplot overlays rescaled observation scores and arrows for the original variables on the first two components [6]. In the course figure, points use the bottom and left scales, and arrows use the top and right scales.




Figure 10. An iris biplot with separate scales for rescaled points and variable arrows.

Arrows pointing the same way mark variables that vary together within the projection. The angle is only as reliable as the two components' representation of each variable. Arrow length and point distance depend on the biplot scaling, so neither is a raw measurement.

R's default biplot uses $\lambda_j = s_j\sqrt{n}$, where s_j is the standard deviation of component j . It divides the scores by λ_j and multiplies the loadings by λ_j before setting the display limits. The Python example uses sample-standardized inputs and the same rescaled quantities, with an explicit arrow factor of 1/60 chosen for display. It shows the quantities without reproducing every R drawing default.

R

```
p <- prcomp(iris[, 1:4], scale. = TRUE)
# R's default biplot rescaling.
lam <- p$sdev[1:2] * sqrt(nrow(iris))
points <- sweep(p$x[, 1:2], 2, lam, "/")
arrows <- sweep(p$rotation[, 1:2], 2, lam, "*")
png("iris-biplot-R.png", width = 700,
    height = 700)
biplot(p)
dev.off()
```

Python

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.decomposition import PCA

features = load_iris(as_frame=True).data
Z = (features - features.mean()) / (
    features.std(ddof=1))
p = PCA().fit(Z)
lam = np.sqrt(
    p.explained_variance_[0:2] * len(Z))
points = p.transform(Z)[0:, 0:2] / lam
arrows = p.components_[0:2].T * lam
fig, ax = plt.subplots(figsize=(6, 6))
ax.scatter(*points.T, s=5)
ax.set(xlim=(-0.26, 0.26), ylim=(-0.26, 0.26),
    xlabel="PC1", ylabel="PC2")
ax.set_aspect("equal")
for (u, v), name in zip(arrows / 60,
    features.columns):
    ax.annotate("", xy=(u, v), xytext=(0, 0),
        arrowprops={"arrowstyle": "->"})
    ax.text(u, v, name, fontsize=7)
# Arrow coordinates use a shared factor 1/60
# for display.
fig.savefig("iris-biplot-Python.png")
plt.close(fig)
```

Optional further reading

Each name below is a link, in the PDF, to the documentation page for that operation. The pages are not additional assessed prerequisites.

- R: scale and Python: StandardScaler explain centering, divisors, and missing values.
- R: cut and Python: qcut document interval boundaries and repeated quantiles.
- R: prcomp and Python: PCA support the optional projection example.

The references below retain the original PCA and biplot papers and the missing-data sources for readers who want further detail.

References

- [1] E. Anderson. “The species problem in Iris”. In: *Annals of the Missouri Botanical Garden* 23.3 (1936), pp. 457–509.
- [2] L. Breiman. “Bagging predictors”. In: *Machine Learning* 24.2 (1996), pp. 123–140.
- [3] S. van Buuren and K. Groothuis-Oudshoorn. “mice: multivariate imputation by chained equations in R”. In: *Journal of Statistical Software* 45.3 (2011), pp. 1–67.
- [4] B. Efron. “Bootstrap methods: another look at the jackknife”. In: *The Annals of Statistics* 7.1 (1979), pp. 1–26.
- [5] R. A. Fisher. “The use of multiple measurements in taxonomic problems”. In: *Annals of Eugenics* 7.2 (1936), pp. 179–188.
- [6] K. R. Gabriel. “The biplot graphic display of matrices with application to principal component analysis”. In: *Biometrika* 58.3 (1971), pp. 453–467.
- [7] D. R. Helsel. “Fabricating data: how substituting values for nondetects can ruin results, and what can be done about it”. In: *Chemosphere* 65.11 (2006), pp. 2434–2439.
- [8] H. Hotelling. “Analysis of a complex of statistical variables into principal components”. In: *Journal of Educational Psychology* 24 (1933), pp. 417–441, 498–520.
- [9] R. J. A. Little and D. B. Rubin. *Statistical Analysis with Missing Data*. 3rd ed. Wiley, 2019.
- [10] K. Pearson. “On lines and planes of closest fit to systems of points in space”. In: *Philosophical Magazine* 2.11 (1901), pp. 559–572.
- [11] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay. “Scikit-learn: machine learning in Python”. In: *Journal of Machine Learning Research* 12 (2011), pp. 2825–2830.
- [12] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. Vienna, Austria, 2026.
- [13] D. B. Rubin. “Inference and missing data”. In: *Biometrika* 63.3 (1976), pp. 581–592.