

Cleaning and Transforming Data

Feature engineering and transparent preparation



Prof. Gabriel Spadon, PhD

Faculty of Computer Science
spadon@dal.ca

Prepare a Table You Can Explain


- Diagnose missing values, variable types, and units.
- Create a derived variable and compare two discretization rules.
- Preserve raw monitoring records and document each cleaning decision.

Missing Data Mechanisms

Mechanism	What missingness can depend on
MCAR	Neither observed nor unobserved values.
MAR	Observed values; after conditioning on them, not the missing values.
MNAR	Unobserved values still matter after conditioning on observed information.

A mechanism describes the recording process. A blank cell alone does not identify why its value is missing.

Alternative Ways to Handle Missing Values



These are alternatives. Keep the original values and record the assumption behind a chosen method.

A Flagged Median Replacement

R

```
raw <- data.frame(value = c(2, NA, 8, 12))
observed <- raw[!is.na(raw$value), ,
               drop = FALSE]
cleaned <- raw
centre <- median(observed$value)
cleaned$was_missing <- is.na(raw$value)
cleaned$value[cleaned$was_missing] <- centre
cleaned
# A chosen value, not a recovered
# measurement.
```


The median of 2, 8, and 12 is 8. The replacement is flagged; the raw table remains available. This rule is not a default for Lab 02.

Python

```
import pandas as pd

raw = pd.DataFrame({"value": [2, None, 8, 12]})
observed = raw.dropna(subset=["value"])
cleaned = raw.copy()
centre = observed["value"].median()
cleaned["was_missing"] = raw["value"].isna()
cleaned["value"] = (
    cleaned["value"].fillna(centre))
print(cleaned)
```

Data Preparation in the Course Workflow



Summaries and plots check preparation decisions. The workflow can loop back.

Four Problems, Four Responses

Variables have incompatible scales → **Data transformation**

A useful ratio needs its own variable → **Feature engineering**

Records contain unavailable or invalid values → **Data cleaning**

Too many variables obscure an analysis → **Dimensionality reduction**

Build a Variable with Meaningful Units

A derived variable combines existing measurements.

$$\text{BMI} = \frac{\text{mass in kg}}{(\text{height in m})^2}$$

For 82 kg and 178 cm, first convert the height to metres.

$$82/1.78^2 = 25.88 \text{ kg/m}^2$$

The unit conversion is part of the calculation. BMI is used here to illustrate arithmetic, not to make a clinical judgement.

Keep the Source Columns

R

```
measurements <- data.frame(  
  mass_kg = c(82, 61, 95),  
  height_cm = c(178, 165, 181))  
measurements$bmi <- measurements$mass_kg /  
  (measurements$height_cm / 100)^2  
round(measurements$bmi, 2)  
# 25.88 22.41 29.00
```

Python

```
import pandas as pd  
  
measurements = pd.DataFrame({  
    "mass_kg": [82, 61, 95],  
    "height_cm": [178, 165, 181]})  
measurements["bmi"] = (  
    measurements["mass_kg"]  
    / (measurements["height_cm"] / 100)**2)  
print(measurements["bmi"].round(2).to_list())  
# [25.88, 22.41, 29.0]
```

Columns from the same table refer to the same rows. R vector arithmetic uses position; pandas Series arithmetic aligns index labels.

Three Different Things You Can Reduce

Fewer variables

Select original features
Project into new
coordinates
Linear (PCA) or nonlinear

Fewer observations

Random sample
Stratified sample
Incremental sample size

Fewer distinct values

Bin numeric values
Combine categories
only with a clear reason

Each reduction discards something. Choose it for an analytical reason, not because the resulting file is smaller.

Random and Stratified Samples

R

```
data <- data.frame(row_id = 1:10,  
  site = rep(c("A", "B"), each = 5))  
set.seed(3141)  
idx <- sample(seq_len(nrow(data)), size = 6)  
random_sample <- data[idx, ]  
# Draw three rows from each site.  
idx_by_site <- split(  
  seq_len(nrow(data)), data$site)  
chosen <- unlist(  
  lapply(idx_by_site, sample, size = 3))  
stratified_sample <- data[chosen, ]  
table(stratified_sample$site)  
# A 3, B 3
```

Python

```
import pandas as pd  
  
data = pd.DataFrame({  
    "row_id": range(1, 11),  
    "site": ["A"]*5 + ["B"]*5})  
random_sample = data.sample(  
    n=6, random_state=3141)  
# Draw three rows from each site.  
stratified_sample = data.groupby("site").sample(  
    n=3, random_state=3141)  
print(stratified_sample["site"].value_counts())  
# A 3, B 3
```

Both samples contain six rows. Stratification fixes three per site. Equal seeds do not imply identical R and Python samples.

Standardize Scale Without Changing Shape

Subtract the mean, then divide by the standard deviation.

$$z_i = \frac{x_i - \bar{x}}{s_x}$$

For $x = 2, 4, 4, 5, 7, 9, 12, 41$, $\bar{x} = 10.5$
and $s_x = 12.72792$.

x_i	$x_i - \bar{x}$	z_i
2	-8.5	-0.668
4	-6.5	-0.511
7	-3.5	-0.275
12	1.5	0.118
41	30.5	2.396

Standardized values have no units. The extreme observation remains extreme; scaling does not remove skewness or make the data normal.

State the Standard Deviation

R

```
x <- c(2, 4, 4, 5, 7, 9, 12, 41)
z <- scale(x)
attr(z, "scaled:center") # 10.5
attr(z, "scaled:scale")  # 12.72792
round(z[1], 4)           # -0.6678
scale(c(x, NA))[1]       # -0.6678231
# Only the last output is NA; observed
# values remain.
```

Python

```
import numpy as np
from sklearn.preprocessing import StandardScaler

x = np.array([2, 4, 4, 5, 7, 9, 12, 41.])
s = StandardScaler().fit(x.reshape(-1, 1))
print(s.mean_[0])           # 10.5
print(round(s.scale_[0], 4)) # 11.9059
print(round(s.transform(x[:, None])[0, 0], 4))
# -0.7139; use ddof=1 for the R convention:
print(round(
    ((x - x.mean()) / x.std(ddof=1))[0], 4))
# -0.6678
```

R uses the sample divisor $n - 1$; StandardScaler uses n . Both fit using observed values and preserve missing positions. A constant column needs an explicit decision before it is scaled.

Boundaries Answer Different Questions

Equal-width

Boundaries are 9.75 units apart.

Interval	Count	Members
[2, 11.75]	6	2,4,4,5,7,9
(11.75, 21.5]	1	12
(21.5, 31.25]	0	none
(31.25, 41]	1	41

Equal-frequency

Boundaries follow the quartiles.

Interval	Count	Members
[2, 4]	3	2, 4, 4
(4, 6]	1	5
(6, 9.75]	2	7, 9
(9.75, 41]	2	12, 41

Both discard within-bin differences. Equal-width retains equal numerical spacing between boundaries; equal-frequency aims for similar counts, subject to ties.

Set Edges, Check Counts

R

```
x <- c(2, 4, 4, 5, 7, 9, 12, 41)
edges <- seq(min(x), max(x), length.out = 5)
width_bins <- cut(x, breaks = edges,
                  include.lowest = TRUE)
quantiles <- quantile(
  x, probs = seq(0, 1, 0.25))
stopifnot(length(unique(quantiles)) == 5)
freq_bins <- cut(x, breaks = quantiles,
                include.lowest = TRUE)
table(width_bins) # 6 1 0 1
table(freq_bins)  # 3 1 2 2
```

Python

```
import numpy as np
import pandas as pd

x = np.array([2, 4, 4, 5, 7, 9, 12, 41.])
edges = np.linspace(x.min(), x.max(), 5)
width_bins = pd.cut(
  x, bins=edges, include_lowest=True)
freq_bins = pd.qcut(x, q=4)
print(pd.Series(width_bins).value_counts(
  sort=False))
print(pd.Series(freq_bins).value_counts(
  sort=False))
```

Include the minimum in the first interval. Repeated quantiles need fewer justified bins or the original numeric variable; do not split identical values across categories.

Encode Categories, Not an Order

R

```
site <- factor(c("A", "B", "A", "C"))
indicators <- model.matrix(~ site - 1)
indicators
# Three columns, one per site; no intercept.
```

Python

```
import pandas as pd

site = pd.Series(["A", "B", "A", "C"])
indicators = pd.get_dummies(site, dtype=int)
print(indicators)
# Three columns, one for each site.
```

Each row has one indicator equal to 1. Coding site A, B, C as 1, 2, 3 would suggest an order and equal spacing that these labels do not have. Lab 02 does not require a model or indicator encoding.

Four Different Value Statuses

Raw field	Status	Interpretation
0	measured	Exact recorded zero in the stated unit.
blank	missing	No usable measurement was recorded.
<0.1	censored	Below 0.1 mg/L; retain 0.1 as a limit, not a value.
oops	invalid	Unexpected token; investigate the original record.
0.24	measured	Exact recorded value, subject to provider quality flags.

Readings are in mg/L. A reporting limit bounds an unobserved measurement; it is not the measurement itself. Keep quality flags and check the provider's definitions.

Preserve the Record, Add an Interpretation

R

```
raw <- data.frame(
  value_text = c("0", "", "<0.1", "oops",
                 "0.24"),
  unit = rep("mg/L", 5),
  stringsAsFactors = FALSE)
cleaned <- raw
cleaned$status <- c(
  "measured", "missing", "censored",
  "invalid", "measured")
cleaned$value_mg_L <- c(0, NA, NA, NA, 0.24)
cleaned$limit_mg_L <- c(NA, NA, 0.1, NA, NA)
stopifnot(identical(
  raw$value_text, cleaned$value_text))
# These statuses follow the stated rules.
# Original strings and units remain.
```

Python

```
import pandas as pd

raw = pd.DataFrame({
  "value_text":
    ["0", "", "<0.1", "oops", "0.24"],
  "unit": ["mg/L"]*5})
cleaned = raw.copy()
cleaned["status"] = [
  "measured", "missing", "censored",
  "invalid", "measured"]
cleaned["value_mg_L"] = [
  0, None, None, None, 0.24]
cleaned["limit_mg_L"] = [
  None, None, 0.1, None, None]
assert raw["value_text"].equals(
  cleaned["value_text"])
```

Validated Conversion, Classify Each Token

R

```
text <- trimws(raw$value_text)
# This provider defines blank as missing.
missing <- text == ""
censored <- grepl("^<[0-9]+([.][0-9]+)?$", text)
measured <- grepl("^[0-9]+([.][0-9]+)?$", text)
cleaned$status <- ifelse(missing, "missing",
  ifelse(censored, "censored",
    ifelse(measured, "measured", "invalid")))
```

Python

```
text = raw["value_text"].str.strip()
# This provider defines blank as missing.
missing = text.eq("")
censored = text.str.fullmatch(
    r"<[0-9]+([.][0-9]+)?")
measured = text.str.fullmatch(
    r"[0-9]+([.][0-9]+)?")
cleaned["status"] = "invalid"
for mask, label in [(missing, "missing"),
                    (censored, "censored"),
                    (measured, "measured")]:
    cleaned.loc[mask, "status"] = label
```

Continue from the raw and cleaned copies. These example rules recognize whole unsigned decimals and less-than limits; unrecognized tokens stay invalid.

Validated Conversion, Convert the Values

R

```
cleaned$value_mg_L <- NA_real_  
cleaned$value_mg_L[measured] <-  
  as.numeric(text[measured])  
cleaned$limit_mg_L <- NA_real_  
cleaned$limit_mg_L[censored] <-  
  as.numeric(sub("^<", "", text[censored]))  
cleaned$value_ug_L <- 1000 * cleaned$value_mg_L
```

Only measured tokens get a value and only censored tokens get a limit; every other row stays missing. The notes explain the patterns and provide the complete script.

Python

```
cleaned["value_mg_L"] = pd.to_numeric(  
    text.where(measured))  
cleaned["limit_mg_L"] = pd.to_numeric(  
    text.where(censored).str.removeprefix("<"))  
cleaned["value_ug_L"] = (  
    1000 * cleaned["value_mg_L"])
```

Check Meaning Before Arithmetic

- Keep the downloaded file unchanged and preserve original strings.
- Identify station, time, parameter, unit, and quality flags; count repeated keys.
- Separate censored limits from exact measurements; investigate invalid tokens.
- Convert compatible units only with a stated rule, such as $1 \text{ mg/L} = 1000 \mu\text{g/L}$.
- Report missing and censored counts beside summaries of usable measurements.

Flagging an unusual reading is a reason to investigate. It is not evidence that the reading should be deleted.

One Record, a Documented Cleaning Path

1. Read one approved monitoring file and identify its observational unit.
2. Audit missing values, reporting limits, types, units, and quality flags.
3. Build a cleaned copy with a reason for each changed or excluded value.
4. Compare summaries before and after; retain the raw record.
5. If useful, compare two binning rules on one suitable numeric variable.

