DALHOUSIE UNIVERSITY

Manipulating and Visualizing Data

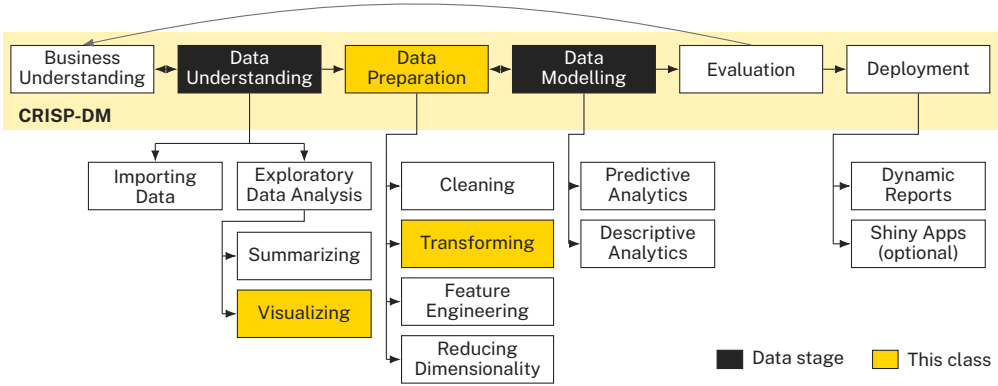
Tables, joins, and graphics in R and Python

Prof. Gabriel Spadon, PhD
Faculty of Computer Science
spadon@dal.ca

From Exploration to a Report

- Today we make the table operations behind a chart or summary explicit.
- Class B combines code, results, and interpretation in a reproducible report.
- Every example is shown in R and Python side by side, on the same records.

Manipulation and Visual Exploration



We move between preparing a useful table and examining what it shows.

Start with the Meaning of a Row

A table operation changes the available evidence.

Our game example has one record per player–game pair. A second table supplies a genre and an achievement total for each listed game.

A Tibble Extends a Data Frame

- A tibble inherits from `data.frame`, but has its own printing and subsetting methods.
- Selecting one column with `tb[, "x"]` keeps a table. Explicit `tb[["x"]]` extracts a vector.
- In pandas, `tb[["x"]]` is a DataFrame and `tb["x"]` is a Series.
- Check a function's documented input behaviour when moving between table classes.

Load the Libraries for Your Language

R

```r
library(dplyr)
library(ggplot2)
```

Python

```python
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
```

Run the examples in order in a working folder where output files can be saved.

Six Operations and Their Effects

`filter`	Keep rows that meet a condition
`select`	Keep or reorder columns
`arrange`	Reorder rows
`mutate`	Add, replace, or remove columns
`group_by`	Define groups without removing rows
`summarise`	Return a summary row per group

Grouping records does not by itself collapse them. A later summary determines the output rows.

Eight Records and a Game Lookup

- Each record is one player and one game.
- Dan has one unknown hours value.
- Tunic appears in the play records but has no lookup entry.
- Stardew appears only in the lookup table.

Keep the original tables. Derived tables will make each decision visible.

Create the Play Records

R

```r
# Player-game records.
plays <- tibble(
  player = c("ana", "ana", "ben", "ben",
             "cara", "cara", "dan", "dan"),
  game = c("Portal", "Celeste", "Portal",
           "Hades", "Celeste", "Hades",
           "Portal", "Tunic"),
  hours = c(12.5, 30, 8, 45.5, 22, 17.5, NA, 6),
  achieved = c(9L, 14L, 4L, 25L, 11L, 8L,
               2L, 3L))
```

Python

```python
# Player-game records.
plays = pd.DataFrame({
    "player": ["ana", "ana", "ben", "ben",
               "cara", "cara", "dan", "dan"],
    "game": ["Portal", "Celeste", "Portal",
             "Hades", "Celeste", "Hades",
             "Portal", "Tunic"],
    "hours": [12.5, 30, 8, 45.5, 22, 17.5,
              np.nan, 6],
    "achieved": [9, 14, 4, 25, 11, 8, 2, 3]})
```

Each row is one player–game record. Hours measure the time played, and achieved counts the achievements earned.

Create the Lookup Table

R

```r
# Game lookup.
games <- tibble(
  game = c("Portal", "Celeste", "Hades",
           "Stardew"),
  genre = c("puzzle", "platformer",
            "roguelike", "farming"),
  total = c(15L, 30L, 49L, 40L))
print(plays)
print(games)
```

Python

```python
# Game lookup.
games = pd.DataFrame({
    "game": ["Portal", "Celeste", "Hades",
             "Stardew"],
    "genre": ["puzzle", "platformer",
              "roguelike", "farming"],
    "total": [15, 30, 49, 40]})
print(plays)
print(games)
```

Continue from the play-record setup. Each game key is unique in this lookup.

Select Records

R

```r
print(filter(plays, hours > 20))
print(filter(plays, hours > 20,
             player != "ana"))
print(filter(plays,
             game %in% c("Hades", "Tunic")))
```

Python

```python
print(plays.loc[plays["hours"] > 20])
print(plays.query(
    "hours > 20 and player != 'ana'"))
print(plays.loc[plays["game"].isin(
    ["Hades", "Tunic"])])
```

The first condition retains Ana–Celeste, Ben–Hades, and Cara–Celeste. Comma-separated conditions in `filter` are combined with AND; pandas keeps the original index labels.

Make the Selection Rule Explicit

- R comparisons with `NA` produce an unknown result. `filter` keeps only true conditions.
- For the floating-point `NaN` in this Python example, `NaN > 20` is false but `NaN != 20` is true.
- Pandas nullable `pd.NA` follows different comparison rules. Do not generalize from one missing-value representation.

Decide whether unknown values belong in the analysis, then include that decision in the mask.

Check the Missing-Value Policy

R

```r
print(nrow(filter(plays, hours > 20)))   # 3
print(nrow(filter(plays, hours != 20)))  # 7
base_plays <- as.data.frame(plays)
print(nrow(
  base_plays[base_plays$hours > 20, ]))  # 4
known <- filter(plays, !is.na(hours),
                hours != 20)
print(nrow(known))  # 7, explicit NA policy
```

Python

```python
print(len(plays.loc[plays["hours"] > 20]))   # 3
print(len(plays.loc[plays["hours"] != 20]))  # 8
known = plays.loc[
    plays["hours"].notna()
    & (plays["hours"] != 20)]
print(len(known))  # 7, explicit NA policy
```

Base bracket selection carries an all-`NA` row for the unknown condition; `filter` keeps only `TRUE`. The explicit known-hours selection gives seven records in both languages.

Select and Rename Variables

R

```r
print(select(plays, player, game, hours))
print(select(plays, -achieved))
print(select(plays, player:hours))
print(select(plays, where(is.numeric)))
print(rename(plays, hrs = hours))
```

Python

```python
print(plays[["player", "game", "hours"]])
print(plays.drop(columns="achieved"))
print(plays.loc[:, "player":"hours"])
print(plays.select_dtypes("number"))
print(plays.rename(columns={"hours": "hrs"}))
```

A label range includes both endpoints. An ordinary Python positional slice excludes its stop position.

Order Records

R

```r
print(arrange(plays, desc(hours)))
print(arrange(plays, player, desc(achieved)))
```

Python

```python
print(plays.sort_values(
    "hours", ascending=False,
    kind="stable", na_position="last"))
print(plays.sort_values(
    ["player", "achieved"],
    ascending=[True, False]))
```

Missing hours appear last. Add ordering columns when ties matter; in pandas, `kind="stable"` keeps the input order of ties.

Derive a Rate

R

```r
rated <- mutate(plays,
  denominator = if_else(hours > 0, hours,
                        NA_real_),
  rate = achieved / denominator,
  band = if_else(rate > 0.5, "hi", "lo"))
print(select(rated, player, game, rate, band))
```

Python

```python
rated = plays.assign(
    denominator=lambda d: d.hours.where(
        d.hours > 0),
    rate=lambda d: d.achieved / d.denominator,
    band=lambda d: d.rate.gt(0.5)
        .map({True: "hi", False: "lo"})
        .where(d.rate.notna()))
print(rated[["player", "game", "rate", "band"]])
```

Rate is achievements per known positive hour. Missing or nonpositive hours produce a missing rate and band.

Report the Summary Denominator

R

```r
print(mean(plays$hours))  # NA
overview <- summarise(plays,
  records = n(),
  players = n_distinct(player),
  observed_hours = sum(!is.na(hours)),
  mean_hours = mean(hours, na.rm = TRUE),
  achievements = sum(achieved))
print(overview)
# 8 records, 4 players, 7 observed,
# 20.214286 h, 76 achievements
```

Python

```python
print(plays["hours"].mean(skipna=False))  # nan
overview = pd.DataFrame([{
    "records": len(plays),
    "players": plays["player"].nunique(),
    "observed_hours": plays["hours"].count(),
    "mean_hours":
        plays["hours"].mean(skipna=True),
    "achievements": plays["achieved"].sum()}])
print(overview)
# 8 records, 4 players, 7 observed,
# 20.214286 h, 76 achievements
```

There are four players, eight play records, and seven observed hours values. The mean describes those seven values.

Read a Pipeline from Top to Bottom

In R, `x |> f()` passes `x` as the first argument of `f`.

A pandas method chain passes the resulting table to the next method. Parentheses allow the chain to span lines.

Choose groups, calculate one summary for each group, then drop the grouping before passing that result on.

Summarize Each Player

R

```r
by_player <- plays |>
  group_by(player) |>
  summarise(records = n(),
    observed_hours = sum(!is.na(hours)),
    known_hours = sum(hours, na.rm = TRUE),
    achievements = sum(achieved),
    .groups = "drop")
print(by_player)
```

Python

```python
by_player = (plays
    .groupby("player", as_index=False)
    .agg(records=("game", "size"),
         observed_hours=("hours", "count"),
         known_hours=("hours", "sum"),
         achievements=("achieved", "sum")))
print(by_player)
```

Dan has six known hours from his Tunic record. His Portal record has unknown hours, so this is a partial observed sum.

Keep the Original Record Count

R

```r
shares <- plays |>
  group_by(player) |>
  mutate(share_known_hours =
    hours / sum(hours, na.rm = TRUE)) |>
  ungroup()
print(select(shares, player, game,
             share_known_hours))
```

Python

```python
shares = plays.assign(share_known_hours=(
    plays.hours
    / plays.groupby("player").hours
      .transform("sum")))
print(shares[["player", "game",
              "share_known_hours"]])
```

Group-wise mutation or transform returns eight records. A summary aggregation instead returns four player rows.

Control Group Order

R

```r
dan_first <- arrange(plays, desc(player))
print(summarise(dan_first,
  achievements = sum(achieved), .by = player))
# First appearance: dan, cara, ben, ana.
```

Python

```python
dan_first = plays.sort_values(
    "player", ascending=False)
print(dan_first.groupby("player", sort=False,
                        as_index=False)
      .agg(achievements=("achieved", "sum")))
# First appearance: dan, cara, ben, ana.
```

The reordered input begins with Dan. First-appearance groups are dan, cara, ben, ana; default grouped summaries sort keys.

Retain an Unobserved Category

R

```r
levels_player <- c("ana", "ben", "cara",
                   "dan", "eve")
p <- mutate(plays,
  player = factor(player,
                  levels = levels_player))
empty_groups <- p |>
  group_by(player, .drop = FALSE) |>
  summarise(records = n(), .groups = "drop")
print(empty_groups)  # eve has 0 recorded rows
```

Python

```python
levels_player = ["ana", "ben", "cara", "dan",
                 "eve"]
p = plays.assign(player=pd.Categorical(
    plays.player, categories=levels_player))
empty_groups = (p
    .groupby("player", observed=False,
             as_index=False)
    .agg(records=("game", "size")))
print(empty_groups)  # eve has 0 recorded rows
```

Eve is a possible category with no recorded rows. Absence from these records does not establish zero real-world activity.

Compose the Six Operations

R

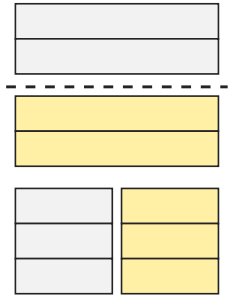
```r
player_rates <- plays |>
  filter(!is.na(hours), hours > 0) |>
  mutate(rate = achieved / hours) |>
  group_by(player) |>
  summarise(records = n(),
            mean_rate = mean(rate),
            .groups = "drop") |>
  arrange(desc(mean_rate)) |>
  select(player, records, mean_rate)
print(player_rates)
```

Python

```python
player_rates = (plays
    .loc[plays.hours.notna()
         & (plays.hours > 0)]
    .assign(rate=lambda d: d.achieved / d.hours)
    .groupby("player", as_index=False)
    .agg(records=("game", "size"),
         mean_rate=("rate", "mean"))
    .sort_values("mean_rate", ascending=False)
    [["player", "records", "mean_rate"]])
print(player_rates)
```

The mean gives every eligible player–game record equal weight. It is not total achievements divided by total hours.

Append Rows or Add Aligned Columns



`bind_rows()`

columns matched by name

`bind_cols()`

rows matched by position

Appending rows aligns column names. Adding columns requires an established correspondence between rows.

Append Compatible Records

R

```r
extra <- tibble(game = "Tunic",
                genre = "puzzle", total = 42L)
extended <- bind_rows(games, extra)
print(extended)  # games itself still has 4 rows
a <- tibble(x = 1)
b <- tibble(y = 2)
print(bind_rows(a, b))  # missing cols become NA
```

Python

```python
extra = pd.DataFrame({"game": ["Tunic"],
    "genre": ["puzzle"], "total": [42]})
extended = pd.concat([games, extra],
                     ignore_index=True)
print(extended)  # games itself still has 4 rows
a = pd.DataFrame({"x": [1]})
b = pd.DataFrame({"y": [2]})
print(pd.concat([a, b], ignore_index=True))
```

Appending does not update the original lookup variable. Check column names, units, and types before binding.

Check Row Correspondence

R

```r
x <- tibble(id = c("a", "b"), value = c(2, 3))
y <- tibble(id = c("a", "b"),
            label = c("low", "high"))
stopifnot(identical(x$id, y$id))
print(bind_cols(x, select(y, label)))
```

Python

```python
x = pd.DataFrame({"id": ["a", "b"],
                  "value": [2, 3]})
y = pd.DataFrame({"id": ["a", "b"],
                  "label": ["low", "high"]})
assert x["id"].equals(y["id"])
print(pd.concat([x, y[["label"]]], axis=1))
```

The identifiers match in the same order. Check this correspondence before attaching columns by position, and do not reset unrelated pandas indexes to force a match.

Track Matched and Unmatched Records

DALHOUSIE UNIVERSITY

7 matched records Portal, Celeste and Hades	1 left-only record Dan's Tunic play	1 right-only game Stardew lookup row

`inner_join` 7 matched records

`left_join` 7 + 1 = 8 records, including Tunic

`right_join` 7 + 1 = 8 rows, including Stardew

`full_join` 7 + 1 + 1 = 9 rows, including both

These counts depend on one lookup row per game. Duplicate lookup keys would create extra matched combinations.

Validate the Lookup Relationship

R

```r
stopifnot(!anyNA(plays$game),
          !anyNA(games$game))
stopifnot(!anyDuplicated(games$game))
joined <- left_join(plays, games, by = "game",
  relationship = "many-to-one",
  na_matches = "never")
stopifnot(nrow(joined) == nrow(plays))
print(joined)
```

Python

```python
assert plays["game"].notna().all()
assert (games["game"].notna().all()
        and games["game"].is_unique)
joined = plays.merge(games, on="game",
                     how="left",
                     validate="many_to_one")
assert len(joined) == len(plays)
print(joined)
```

Complete keys and a unique lookup key are checked before the merge. The left join preserves all eight play records.

Compare the Four Join Counts

R

```r
print(nrow(joined))                    # 8
print(nrow(inner_join(plays, games,
                      by = "game")))   # 7
print(nrow(right_join(plays, games,
                      by = "game")))   # 8
print(nrow(full_join(plays, games,
                     by = "game")))    # 9
```

Python

```python
print(len(joined))                     # 8
print(len(plays.merge(games, on="game",
                      how="inner")))   # 7
print(len(plays.merge(games, on="game",
                      how="right")))   # 8
print(len(plays.merge(games, on="game",
                      how="outer")))   # 9
```

Inner has seven rows; left and right each have eight; full or outer has nine.

Make Cardinality an Explicit Expectation

Check	dplyr	pandas
Lookup cardinality	`relationship = "many-to-one"`	`validate="many_to_one"`
Missing keys	May match by default; `na_matches="never"` prevents this	May match by default; check or separate missing keys before merging
Output order	Left rows first, then unmatched right rows	Outer merge sorts keys; sort explicitly if order matters
Name clashes	Default suffixes `.x`, `.y`	Default suffixes `_x`, `_y`

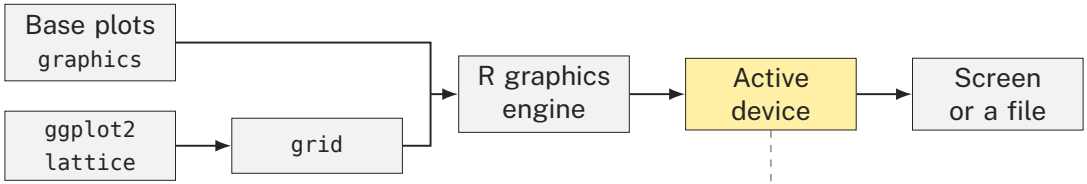
Validate the lookup relationship and inspect unmatched records.

Draw from a Checked Table

The figure should answer a question already defined by the table and its eligibility rules.

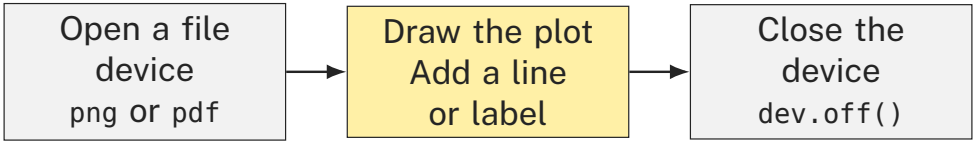
We revisit the Iris measurements from Week 3 to compare plotting interfaces, then return to the game summary.

From Drawing Commands to an Output



Base graphics and grid use the graphics engine and an active output device.

Open, Draw, and Close



Drawing commands go to the active device.
Closing the file device completes that output.

Interactive sessions may draw on screen; a script may choose another default.
Select the intended file device explicitly.

Load the Iris Measurements

R

```r
d <- as_tibble(iris)
print(nrow(d))                  # 150 flowers
print(mean(d$Sepal.Width))      # 3.057333 cm
```

Python

```python
raw = load_iris(as_frame=True)
d = raw.data.copy()
d.columns = ["Sepal.Length", "Sepal.Width",
             "Petal.Length", "Petal.Width"]
d["Species"] = pd.Categorical.from_codes(
    raw.target, raw.target_names)
print(len(d), d["Sepal.Width"].mean())
# 150, 3.057333 cm
```

The same 150 flower measurements and centimetre units used in Week 3 are available locally in both languages.

Save the Scatter Plot

R

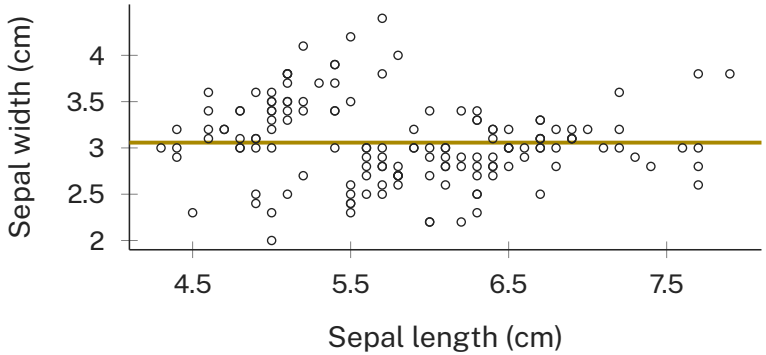
```r
png("class07-iris-base.png", width = 1200,
    height = 800, res = 150)
plot(d$Sepal.Length, d$Sepal.Width, pch = 1,
     xlab = "Sepal length (cm)",
     ylab = "Sepal width (cm)")
abline(h = mean(d$Sepal.Width),
       col = "#A88800", lwd = 2)
dev.off()
```

Python

```python
fig, ax = plt.subplots(figsize=(8, 16 / 3))
ax.scatter(d["Sepal.Length"], d["Sepal.Width"],
           facecolors="none",
           edgecolors="black")
ax.axhline(d["Sepal.Width"].mean(),
           color="#A88800", linewidth=2)
ax.set(xlabel="Sepal length (cm)",
       ylabel="Sepal width (cm)")
fig.savefig("class07-iris-base.png", dpi=150)
plt.close(fig)
```

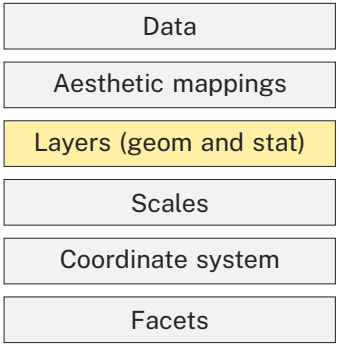
The gold line marks the overall mean sepal width, 3.057333 cm. The PNG is 1200 by 800 pixels.

Show the Overall Mean as a Reference



Measured to 0.1 cm, the 150 flowers share 117 positions, so some overlap. The gold line is the mean width, not a fitted relationship.

Separate the Choices That Make a Plot



A point layer displays the original values.

A histogram layer counts values in bins.

Defaults supply many settings, but data and mappings alone draw no marks.

The grammar describes how data, marks, calculations, and display settings work together. See Wilkinson and the ggplot2 book in the notes.

Connect Each Element to the Iris Plot

DALHOUSIE UNIVERSITY

Element	Meaning	Iris example
Data	Values to display	Flower measurements
Aesthetic mappings	Variables linked to visible properties	Length to x, width to y
Geometric objects	Marks that display values	One point per flower
Statistical transformations	Calculations used by a layer	Identity for points; counts for histogram bins

A layer combines marks with a calculation.

Choose How the Values Are Displayed

Element	Meaning	Iris example
Scales	Convert values into positions, colours, or other properties	Numeric axes and species colours
Coordinate system	Geometry of the plotting panel	Cartesian x and y
Faceting	Separate panels for subsets	One panel per species

A scale need not be one-to-one. Several values can share a binned colour category.

Draw One Panel per Species

R

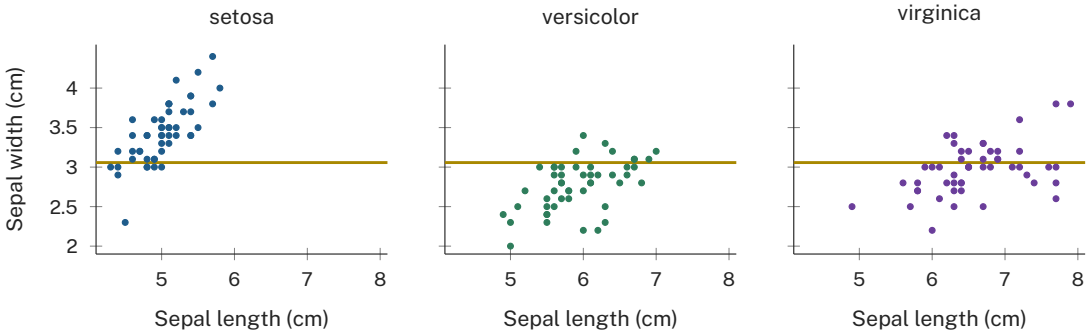
```r
colours <- c("#1F5C8B", "#2E7D5B", "#6A3D9A")
p_iris <- ggplot(d,
  aes(Sepal.Length, Sepal.Width,
      colour = Species)) +
  geom_point() +
  geom_hline(yintercept = mean(d$Sepal.Width),
             colour = "#A88800") +
  facet_wrap(~ Species) +
  scale_colour_manual(values = colours) +
  labs(x = "Sepal length (cm)",
       y = "Sepal width (cm)") +
  theme_minimal()
ggsave("class07-iris-facets.pdf", p_iris,
       width = 9, height = 3.5)
```

Python

```python
colours = ["#1F5C8B", "#2E7D5B", "#6A3D9A"]
sns.set_theme(style="whitegrid")
g = sns.relplot(
    data=d, x="Sepal.Length", y="Sepal.Width",
    hue="Species", col="Species",
    kind="scatter", palette=colours,
    height=3.5, aspect=6 / 7)
g.refline(y=d["Sepal.Width"].mean(),
          color="#A88800")
g.set_axis_labels("Sepal length (cm)",
                  "Sepal width (cm)")
g.savefig("class07-iris-facets.pdf")
plt.close(g.figure)
```

Each species has its own labelled panel. Every panel uses the overall width mean as the same gold reference.

Compare the Species Panel by Panel



The panels share both axes. Of 50 flowers per species, 42 setosa, 8 versicolor and 17 virginica lie above the gold mean width.

Check Eligibility Before Summarizing

- Return to the play records and their checked left join.
- This question concerns records with known hours greater than five and a known positive achievement total.
- Retain the two excluded rows for inspection. Do not let a grouping default decide which evidence disappears.
- Summarize the six eligible records, then draw that summary.

Retain an Exclusion Record

R

```
checked <- mutate(joined,
  eligible = !is.na(hours) & hours > 5 &
    !is.na(total) & total > 0 & !is.na(genre))
excluded <- filter(checked, !eligible)
eligible <- filter(checked, eligible)
print(select(excluded, player, game, hours,
             genre, total))
print(c(input = nrow(plays),
        eligible = nrow(eligible),
        excluded = nrow(excluded)))  # 8, 6, 2
```

Python

```
checked = joined.assign(eligible=
    joined.hours.notna() & (joined.hours > 5) &
    joined.total.notna() & (joined.total > 0) &
    joined.genre.notna())
excluded = checked.loc[~checked.eligible].copy()
eligible = checked.loc[checked.eligible].copy()
print(excluded[["player", "game", "hours",
                "genre", "total"]])
print(len(plays), len(eligible),
      len(excluded))  # 8, 6, 2
```

Dan–Portal is excluded because hours are unknown. Dan–Tunic is excluded because its game lookup and denominator are missing.

Compute the Genre Summary

R

```r
genre_summary <- eligible |>
  mutate(pct = 100 * achieved / total) |>
  group_by(genre) |>
  summarise(records = n(),
            known_hours = sum(hours),
            pct = mean(pct),
            .groups = "drop") |>
  arrange(desc(pct))
print(genre_summary)
```

Python

```python
genre_summary = (eligible
    .assign(pct=lambda d:
            100 * d.achieved / d.total)
    .groupby("genre", as_index=False)
    .agg(records=("player", "size"),
         known_hours=("hours", "sum"),
         pct=("pct", "mean"))
    .sort_values("pct", ascending=False))
print(genre_summary)
```

Each eligible play record receives equal weight. All three genres have two records.

Save the Genre Figure

R

```r
p_genre <- ggplot(genre_summary,
  aes(reorder(genre, -pct), pct)) +
  geom_col(fill = "#1F5C8B",
           colour = "#242424") +
  geom_text(aes(label = round(pct, 1)),
            vjust = -0.4) +
  scale_y_continuous(limits = c(0, 50)) +
  labs(x = "Genre",
       y = "Mean achievement completion (%)") +
  theme_minimal()
ggsave("class07-genre-summary.pdf", p_genre,
       width = 7, height = 4)
```

Python

```python
fig, ax = plt.subplots(figsize=(7, 4))
bars = ax.bar(genre_summary.genre,
              genre_summary.pct,
              color="#1F5C8B",
              edgecolor="#242424")
ax.bar_label(bars, fmt="%.1f")
ax.set(xlabel="Genre",
       ylabel="Mean achievement completion (%)",
       ylim=(0, 50))
fig.tight_layout()
fig.savefig("class07-genre-summary.pdf")
plt.close(fig)
```

Both languages draw the same ordered bar chart with its labelled values and save it as a PDF.

Save the Summary and the Exclusions

R

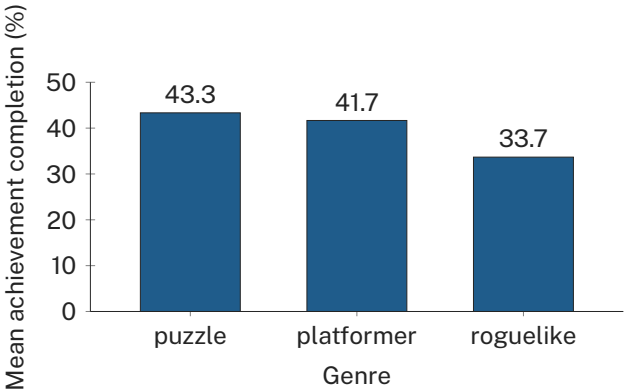
```r
write.csv(genre_summary,
          "class07-genre-summary.csv",
          row.names = FALSE)
write.csv(excluded, "class07-excluded.csv",
          row.names = FALSE)
```

Python

```python
genre_summary.to_csv(
    "class07-genre-summary.csv", index=False)
excluded.to_csv("class07-excluded.csv",
                index=False)
```

The saved CSV, chart, and exclusion record provide inputs for a report. Keep the source tables and the code as well.

Read the Mean Completion Percentage



Puzzle leads platformer by 1.7 percentage points; roguelike trails it by 8.0.

Make Every Transformation Explainable

- Say what each row represents and how many values a summary uses.
- Keep original data and retain a record of exclusions.
- Validate lookup keys and distinguish unmatched rows from duplicated matches.
- Save a labelled figure with its numerical summary and the code that produced both.