

Reproducible Reporting

R Markdown and Python notebooks


Prof. Gabriel Spadon, PhD


Faculty of Computer Science
spadon@dal.ca

Turn an Analysis into a Reproducible Report

- Combine a question, code, evidence and interpretation.
- Produce a PDF from a fresh R process or Python kernel.
- Distinguish execution errors from conversion errors.
- Record the inputs, parameters and environment.


Choose R Markdown or a Python notebook. The complete examples in the notes support either route.

Communicate the Analysis to Its Reader



Dynamic reports communicate earlier work; Shiny apps are an optional sibling.


A Copied Result Can Become Stale



Copied numbers can remain stale after the analysis changes.
Run the linked report again and inspect the output.

Computed text can update a count after a rerun. An ordinary interpretation sentence still needs to be checked.


Connect Explanation and Computation



A source links the analysis to the report.
Run it again to update the output.

A PDF records one run. The source still needs its input data and software to reproduce that run.


Separate Execution from Conversion



R and knitr execute enabled chunks. Pandoc converts the document; XeLaTeX compiles the PDF.

Execute Cells Before Exporting the Report

--execute runs
cells in order
before conversion.



An ipynb file stores cells, metadata and possibly old outputs. Exporting those saved outputs alone does not recompute the analysis.

Keep the Source and Inputs Together

For Lab 04, work in your own repository's `labs/lab-04` folder.

- Save `report.Rmd` or `report.ipynb`.
- Keep input paths relative to that working folder.
- Retain raw data and record preparation choices.
- Run the command in the folder that contains the report.

Use a separate practice folder for each class example. The notes contain all source text and the full synthetic CSV.

Prepare One Language Environment

R Markdown

R and an editor such as RStudio.

R packages `rmarkdown`, `knitr`, `ggplot2`, `dplyr`.

Pandoc and XeLaTeX for the PDF route.

Python notebook

JupyterLab and a Python kernel.

`pandas`, `NumPy`, `matplotlib`, `scikit-learn`, `IPython`, `nbconvert` and `Jinja2`.

Pandoc and XeLaTeX for the PDF route.

Install dependencies once in the chosen environment. The notes give the setup and kernel-selection commands.

Save Files Beyond the Runtime

- Local RStudio or Jupyter keeps files on your computer.
- Posit Cloud provides projects that can be downloaded.
- Colab saves notebooks separately from its temporary runtime files.
- For Kaggle or another supported option, check language, persistence and export before starting.

If cloud PDF tools are unavailable, download the source and inputs and render locally. Keep coursework private until the deadline and instructor release.

Use the Familiar Iris Measurements

Iris contains 150 flowers, 50 of each species, measured in centimetres.

- One compact table compares species means.
- A histogram shows the pooled sepal-length distribution.
- Box plots compare variation within and between species.

The R and scikit-learn copies agree. This historical example does not represent a random sample of every iris population.

Start with Document Settings

```
---  
title: "Iris sepal lengths"  
author: "Course example"  
date: ""  
output:  
  pdf_document:  
    latex_engine: xelatex  
    fig_caption: true  
  html_document: default  
---
```

The YAML header names the output formats. Its first format is the default; a render command can select another.

Run Setup Without Printing It

```
```{r setup, include=FALSE}  
library(ggplot2)
library(dplyr)
knitr::opts_chunk$set(
 echo = FALSE, cache = FALSE, error = FALSE,
 warning = TRUE, message = FALSE,
 fig.width = 6, fig.height = 3, out.width = "90%"
)
dat <- iris
```
```

The setup executes. `include=FALSE` hides its code and output, while `cache=FALSE` keeps this small analysis fresh.

Put a Calculation into an R Sentence

Question and data

How do sepal lengths vary within and between the three species in this sample? Iris contains 150 flowers, 50 of each species, with measurements in centimetres. These historical example data are not a random sample of all iris populations.

The mean sepal length is ``r sprintf("%.2f", mean(dat$Sepal.Length))`` cm, with sample standard deviation ``r sprintf("%.3f", sd(dat$Sepal.Length))`` cm.

A compact summary

The rendered mean is 5.84 cm and sample standard deviation is 0.828 cm. These values update when this source is rerendered.

Emit the Same Sentence from Python

Run the Iris Python setup cell in the notes before this example.

```
mean_sepal = dat["Sepal_cm"].mean()
sd_sepal = dat["Sepal_cm"].std(ddof=1)
display(
    Markdown(
        f"The mean sepal length is {mean_sepal:.2f} cm, "
        f"with sample standard deviation {sd_sepal:.3f} cm."
    )
)
```

A normal Markdown cell does not evaluate Python. This code cell formats its result as Markdown output; ddof=1 matches R's sample standard deviation.

Compute a Compact Group Summary

R

```
```{r summary}
summary_table <- dat |>
 group_by(Species) |>
 summarise(
 n = n(), Sepal_cm = mean(Sepal.Length),
 Petal_cm = mean(Petal.Length), .groups = "drop"
)
knitr::kable(summary_table, row.names = FALSE, digits = 3)
```
```

Python

```
summary_table = (
    dat.groupby("Species", sort=True)
    .agg(
        n=("Sepal_cm", "size"),
        Sepal_cm=("Sepal_cm", "mean"),
        Petal_cm=("Petal_cm", "mean"),
    )
    .reset_index()
)
```

Both give one row per species with a count and two means. In R, kable formats the table; Python keeps Species as a named column.

Keep Counts and Units Visible

| Species | n | Sepal mean (cm) | Petal mean (cm) |
|------------|-----|-----------------|-----------------|
| setosa | 50 | 5.006 | 1.462 |
| versicolor | 50 | 5.936 | 4.260 |
| virginica | 50 | 6.588 | 5.552 |

Every species has 50 observations. Means alone do not reveal the spread, so the report also needs plots.

Supply the Table to Both Export Formats

```
display(  
  {  
    "text/html": summary_table.to_html(  
      index=False, float_format="%.3f", escape=True  
    ),  
    "text/latex": summary_table.to_latex(  
      index=False, float_format="%.3f", escape=True  
    ),  
  },  
  raw=True,  
)
```


HTML and LaTeX are two representations of the same calculation. Each exporter selects the one it can use.

Record the Bin Grid in the Report

```
```{r histogram, fig.cap="Sepal length distribution."}  
edges <- seq(4.25, 8.25, length.out = 11)
ggplot(dat, aes(Sepal.Length)) +
 geom_histogram(
 breaks = edges, closed = "left",
 fill = "#D4A017", colour = "white"
) +
 labs(x = "Sepal length (cm)", y = "Number of flowers") +
 theme_minimal(base_size = 11)
```
```


The same Week 3 grid is used in Python. The ten counts sum to 150; none of the measurements lies on an internal edge.

Read the Pooled Distribution



Bins have width 0.4 cm and run from 4.25 to 8.25 cm. The plot combines all three species.

Compare the Species with Box Plots



Both report routes use linear quartiles and 1.5-IQR whiskers, as in Week 3. The separate virginica point is 4.9 cm; its lower whisker is 5.6 cm.

Distinguish Device Size from Display Size

R

`fig.width` and `fig.height` are in inches.

`fig.asp` is a height-to-width ratio.

`out.width` controls display width in the report.

Python

`figsize=(6,3)` sets a six-by-three inch figure.

Label the axes, apply `tight_layout()`, then check the exported size.

Shrinking a figure after drawing it also shrinks the labels. A caption should explain the units and plotting convention.

Render from a Fresh Process or Kernel

```
Rscript --vanilla -e "rmarkdown::render(  
  'report.Rmd', output_format='pdf_document')"
```

```
jupyter nbconvert --execute --to pdf --no-input report.ipynb
```

Run the command for your language in the report folder. Both produce `report.pdf`. The notebook flag `--no-input` hides code; `--execute` still runs it.

Check the Result That Will Be Submitted

1. Confirm the command completed without an execution error.
2. Open the new PDF and inspect every page.
3. Verify one count and calculated value against the input.
4. Check every plot, unit, table and interpretation.
5. Preserve the source, input identity and software versions.

A successful render proves this configured run completed. It does not prove that the analysis or conclusion is correct.

Separate Running Code from Showing Code

| Option | R Markdown default | Effect |
|---------|--------------------|---------------------------------|
| eval | TRUE | Run the code |
| echo | TRUE | Show source code |
| include | TRUE | Include source and output |
| results | 'markup' | Format text output |
| message | TRUE | Include messages |
| warning | TRUE | Include warnings |
| error | FALSE | Stop on a code error |
| cache | FALSE | Reuse cached results if enabled |

Compare Eval, Echo and Include

- `eval=FALSE` disables code execution.
- `echo=FALSE` hides the source code.
- `include=FALSE` omits source and output.

Use a unique label such as `species-summary`. Hyphens are a safe convention, though a period is not prohibited. The required analysis must still execute when the report renders.

Use the Settings of the Selected Output

`fig.cap` supplies an R figure caption, and the example's PDF settings enable numbered figure captions.

- Numbering and float placement depend on the output format.
- Python examples put a numbered title inside each figure.
- Keep captions beside their figures in the final PDF.
- HTML is useful for checking, but Lab 04 requires PDF.

The notes also explain results, message and error options, and optional Word, Markdown and Beamer outputs.

Understand What a Cache Reuses

The small class reports `keep cache=FALSE`.

- A cache can skip code whose prior result appears reusable.
- knitr tracks code and relevant options, not every external file change.
- A file checksum can be an explicit cache dependency.
- Setup and side effects need particular care.


The notes show a cached descriptive summary. No model or cache is required for the brief. A fresh notebook execution recomputes its code cells.

Saved Output Can Come from an Earlier Run

- R notebooks support interactive chunk output and HTML previews.
- Python notebooks can retain outputs from cells run out of order.
- An output-format name does not itself start a fresh session.
- Rebuild with the tested process or kernel before final export.

An ipynb is JSON text. Its outputs and metadata can make diffs noisy, but its source and history still belong in the reproducible record.

Name an Input Choice and Record It



A parameter changes the selection while preserving one analytical recipe. The report should state the value used.

Use Six Synthetic Game Records

game,players,completed,achievements

Celeste,412,254,32

Hades,388,180,49

Stardew,301,112,40

Factorio,154,39,27

Spelunky,88,21,25

Tunic,64,12,26

All values are fictional teaching counts. This is a separate game-level dataset, not an aggregation of Class A's play records. Save the CSV beside the report.

Define the Completion Rate First

$$\text{completion rate} = \frac{\text{completed players}}{\text{listed players}}$$

- A row represents one game.
- Completed players cannot exceed listed players.
- Players must be positive for a defined rate.
- Achievement count is not the rate denominator.

The full source checks missing, negative, non-integer and inconsistent counts before creating a filtered copy.

Declare the Default Values

R

```
---  
title: "Completion rates in a synthetic game dataset"  
author: "Course example"  
date: ""  
output:  
  pdf_document:  
    latex_engine: xelatex  
    fig_caption: true  
  html_document: default  
params:  
  datafile: "steam_achievements.csv"  
  minplayers: 100  
---
```

Both declare the same two defaults. The datafile names a sibling CSV with this schema; Iris would not fit it.

Python

```
datafile = "steam_achievements.csv"  
minplayers = 100
```



A plain cell near the top holds the defaults. The preparation cell after it validates the counts, filters a copy of raw and computes the rate.

Verify the Default Table

| Game | Players | Completed | Achievements | Rate |
|----------|---------|-----------|--------------|-------|
| Celeste | 412 | 254 | 32 | 0.617 |
| Hades | 388 | 180 | 49 | 0.464 |
| Stardew | 301 | 112 | 40 | 0.372 |
| Factorio | 154 | 39 | 27 | 0.253 |

At least 100 players retains four of six games. Completed divided by players gives each rate; the achievement count remains a separate quantity.

Show Rates and Their Denominators



Keep games in the same order. The rate plot compares proportions; the count plot reveals how many players each proportion describes.

Preserve a Distinct Output for Each Choice

```
Rscript --vanilla -e "rmarkdown::render(  
  'report.Rmd', params=list(minplayers=300),  
  output_file='report-300.pdf')"
```

```
jupyter nbconvert --execute --to pdf --no-input --output report-300 report.ipynb
```

For Python, first change and save the parameter cell to 300. Both versions retain three games. Their per-game rates stay the same.

Report an Empty Selection Honestly

With `minplayers = 1000`, none of the six games is eligible.

- State that no games met the threshold.
- Do not present an empty chart as a meaningful comparison.
- Preserve the chosen threshold in the report.

Changing a filter changes coverage. These synthetic values support no claims about real Steam users or games.

Run the Report for Several Values

R can pass `params=list(...)` to render; RStudio also offers a parameter dialog.

Python can use Papermill with a cell tagged parameters.

Keep distinct output names and record each parameter value.

The notes give both batch examples. The ordinary edit-and-rerun workflow is sufficient for Lab 04.

Read the First Informative Error

- Missing name or package: check setup and the selected environment.
- Missing file: check the working folder and relative path.
- Pandoc error: check document conversion.
- XeLaTeX error: check the PDF compilation log.


A working HTML export can help isolate a PDF-stage problem. It does not prove that the final PDF is correct or that every required computation ran.

Preserve More Than the Final PDF

- Source, raw input identity, parameters and software versions.
- Units, exclusions, cleaning decisions and analytical limits.
- The PDF from the final fresh run.
- Meaningful commits and the submitted full commit SHA.

Preserve existing history. Do not amend, squash, rebase, force-push or backdate submitted work. Publish only after the deadline and instructor release.

Several Tools Share the Same Reporting Idea



Literate programming connects explanation and code. Historical tools and Quarto are context, not extra environments to install this week.

Use the Reference for Your Chosen Route

- [R Markdown: The Definitive Guide](#) and the [render reference](#) explain document formats, execution and parameters.
- The [knitr options reference](#) and [Posit cheatsheets](#) support the R examples.
- The [nbconvert user guide](#) and [IPython display reference](#) support notebook execution, PDF export and computed prose.
- For extensions, see [Papermill parameters](#), [Quarto](#), and the [Pandoc manual](#).

Build, Change and Explain One Report

1. Choose R or Python and build the Iris report.
2. Export a fresh PDF and verify its count, mean and plots.
3. Change the synthetic game threshold from 100 to 300.
4. Explain the coverage change and distinguish rate from count.

The small worked examples are ungraded practice. The Lab 04 handout defines the assessed dataset, language, and submission package.

Lab 04 | Sustainable communities brief

Turn one verified public dataset into a concise reproducible brief.

Carry forward

Data audit, cleaning, table manipulation, graphics, and reporting (Weeks 01–04).

Start in class

Choose one verified dataset. Work through the selection checklist, then set up the R or Python analysis.

One assessment, two languages

Submit either R or Python. Both paths use the same evidence and rubric.

Handout

`week-04/Assessment/lab-04`